

5. Online Speech-to-Text (ASR)

5. Online Speech-to-Text (ASR)

1. Concept Introduction

1.1 What is "ASR"?

1.2 Implementation Principles

1. Acoustic Model

2. Language Model

3. Pronunciation Dictionary

4. Decoder

5. End-to-End ASR

2. Project Architecture

Key Code

1. Speech Processing and Recognition Core (`targetmodel/targetmodel/asr.py`)

2. VAD Smart Recording (`targetmodel/targetmodel/asr.py`)

Code Analysis

3. Practical Operations

3.1 Configuring Online ASR

3.2 Startup and Testing

1. Concept Introduction

1.1 What is "ASR"?

ASR (Automatic Speech Recognition) is a technology that converts human speech signals into text. It is widely used in intelligent assistants, voice command control, automated telephone customer service, real-time caption generation, and other fields. The goal of ASR is to enable machines to "understand" human language and convert it into a form that computers can process and understand.

1.2 Implementation Principles

The implementation of an ASR system mainly relies on the following key technical components:

1. Acoustic Model

- The acoustic model is responsible for converting the input sound signal into phonemes or sub-word units. This typically involves feature extraction steps, such as using Mel-frequency cepstral coefficients (MFCCs) or filter banks to represent the audio signal.
- These features are then fed into deep neural networks (DNNs), convolutional neural networks (CNNs), recurrent neural networks (RNNs), or more advanced Transformer architectures for training to learn how to map audio features to corresponding phonemes or characters.

2. Language Model

- A language model is used to predict the next most likely word given context, thus helping to improve recognition accuracy. It is trained on large amounts of text data to understand which word sequences are more likely to occur.
- Common language models include n-gram models, RNN-based LMs, and the recently popular Transformer-based LMs.

3. Pronunciation Dictionary

- A pronunciation dictionary provides a mapping between words and their corresponding pronunciations. This is crucial for connecting the acoustic model and the language model, as it allows the system to understand and match heard sounds based on known pronunciation rules.

4. Decoder

- The decoder's task is to find the most likely word sequence as output, given the acoustic model, the language model, and the pronunciation dictionary. This process typically involves complex search algorithms, such as the Viterbi algorithm or graph-based search methods, to find the optimal path.

5. End-to-End ASR

- With the development of deep learning, end-to-end ASR systems have emerged. These systems attempt to learn text output directly from raw audio signals without requiring explicit separate designs for acoustic models, pronunciation dictionaries, and language models. These systems are often based on sequence-to-sequence (Seq2Seq) frameworks, such as using attention mechanisms or Transformer architectures, greatly simplifying the complexity of traditional ASR systems.

In summary, modern ASR systems achieve efficient and accurate human speech-to-text conversion by combining the above components and training with large-scale datasets and powerful computing resources. With technological advancements, the performance of ASR systems continues to improve, and their application scenarios are becoming increasingly widespread.

2. Project Architecture

Key Code

1. Speech Processing and Recognition Core (`largemodel/largemodel/asr.py`)

```
# From largemodel/largemodel/asr.py
def kws_handler(self)->None:
    if self.stop_event.is_set():
        return

    if self.listen_for_speech(self.mic_index):
        asr_text = self.ASR_conversion(self.user_speechdir) # Perform ASR
        conversion
        if asr_text == 'error': # Check if ASR result length is less than 4
            characters
            self.get_logger().warn("I still don't understand what you mean.
            Please try again")
            playsound(self.audio_dict[self.error_response]) # Error response
        else:
            self.get_logger().info(asr_text)
            self.get_logger().info("okay 🤖, let me think for a moment...")
            self.asr_pub_result(asr_text) # Publish ASR result
    else:
        return
```

```

def ASR_conversion(self, input_file:str)->str:
    if self.use_online_asr:
        result=self.modelinterface.online_asr(input_file)
        if result[0] == 'ok' and len(result[1]) > 4:
            return result[1]
        else:
            self.get_logger().error(f'ASR Error:{result[1]}') # ASR error.
            return 'error'
    else:
        result=self.modelinterface.SenseVoiceSmall_ASR(input_file)
        if result[0] == 'ok' and len(result[1]) > 4:
            return result[1]
        else:
            self.get_logger().error(f'ASR Error:{result[1]}') # ASR error.
            return 'error'

```

2. VAD Smart Recording (largemodel/largemodel/asr.py)

```

# From largemodel/largemodel/asr.py
def listen_for_speech(self,mic_index=0):
    p = pyaudio.PyAudio() # Create PyAudio instance.
    audio_buffer = [] # Store audio data.
    silence_counter = 0 # Silence counter.
    MAX_SILENCE_FRAMES = 90 # Stop after 900ms of silence (30 frames * 30ms)
    speaking = False # Flag indicating speech activity.
    frame_counter = 0 # Frame counter.
    stream_kwargs = {
        'format': pyaudio.paInt16,
        'channels': 1,
        'rate': self.sample_rate,
        'input': True,
        'frames_per_buffer': self.frame_bytes,
    }
    if mic_index != 0:
        stream_kwargs['input_device_index'] = mic_index

    # Prompt the user to speak via the buzzer.
    self.pub_beep.publish(UInt16(data = 1))
    time.sleep(0.5)
    self.pub_beep.publish(UInt16(data = 0))

    try:
        # Open audio stream.
        stream = p.open(**stream_kwargs)
        while True:
            if self.stop_event.is_set():
                return False

            frame = stream.read(self.frame_bytes, exception_on_overflow=False)
            # Read audio data.
            is_speech = self.vad.is_speech(frame, self.sample_rate) # VAD
            detection.

            if is_speech:
                # Detected speech activity.
                speaking = True
                audio_buffer.append(frame)

```

```

        silence_counter = 0
    else:
        if speaking:
            # Detect silence after speech activity.
            silence_counter += 1
            audio_buffer.append(frame) # Continue recording buffer.

            # End recording when silence duration meets the threshold.
            if silence_counter >= MAX_SILENCE_FRAMES:
                break
        frame_counter += 1
        if frame_counter % 2 == 0:
            self.get_logger().info('1' if is_speech else '-')
            # Real-time status display.
    finally:
        stream.stop_stream()
        stream.close()
        p.terminate()

    # Save valid recording (remove trailing silence).
    if speaking and len(audio_buffer) > 0:
        # Trim the last silent part.
        clean_buffer = audio_buffer[:-MAX_SILENCE_FRAMES] if len(audio_buffer) >
MAX_SILENCE_FRAMES else audio_buffer

        with wave.open(self.user_speechdir, 'wb') as wf:
            wf.setnchannels(1)
            wf.setsampwidth(p.get_sample_size(pyaudio.paInt16))
            wf.setframerate(self.sample_rate)
            wf.writeframes(b''.join(clean_buffer))
        return True

```

Code Analysis

The ASR (Speech-to-Text) function is provided by the `ASRNode` node (`asr.py`). This node is responsible for recording, converting, and publishing audio.

1. Audio Recording (`listen_for_speech`):

- This function uses the `pyaudio` library to capture an audio stream from the microphone.
- It integrates the `webrtcvad` library for Voice Activity Detection (VAD). The function loops through audio frames and uses `vad.is_speech()` to determine if each frame contains human voice.
- When speech is detected, data is written to a buffer. Recording stops when sustained silence (defined by `MAX_SILENCE_FRAMES`) is detected.
- Finally, the audio data in the buffer is written to a `.wav` file at `self.user_speechdir`.

2. Backend Selection and Execution (`ASR_conversion`):

- The `kws_handler` function calls the `ASR_conversion` function after successful recording.
- This function determines which backend implementation to call by reading the ROS parameter `use_online_asr` (a boolean value).
- If `false`, the `self.modelinterface.Sensevoicesmall_ASR` method is called, corresponding to local recognition.

- If `true`, the `self.modelinterface.oline_asr` method is called, corresponding to online recognition.
- This function takes the audio file path as an argument to the selected method and processes the returned result.

3. Result Publication (`asr_pub_result`):

- After `ASR_conversion` returns valid text, `kws_handler` calls the `asr_pub_result` function.
- This function wraps the text string in a `std_msgs.msg.String` message and publishes it to the `/asr` topic via the ROS publisher.

3. Practical Operations

3.1 Configuring Online ASR

1. First obtain an API Key from the Tongyi Qianwen registration tutorial
2. Open the model interface configuration file `large_model_interface.yaml`:

```
vim
~/yahboom_ws/install/largemodel/share/largemodel/config/large_model_interfac
e.yaml
```

3. Enter your API credentials:

Find the online ASR section, enter your Key, and select the speech recognition model.

```
# Large Language Models Settings
# -----
---
## Online Large Language Models
# Tongyi Platform Configuration
tongyi_api_key: "Your_AccessKey_ID" # API Key

# -----
---
# Speech Recognition Settings
# -----
---
## Online ASR
oline_asr_model: 'paraformer-realtime-8k-v2' # Options: 'paraformer-
realtime-v2', 'paraformer-realtime-v1', etc., uses Tongyi Qianwen platform
API
oline_asr_sample_rate: 16000 # ASR audio sample rate
```

4. Open the main configuration file `yahboom.yaml`:

```
vim ~/yahboom_ws/install/largemodel/share/largemodel/config/yahboom.yaml
```

5. Enable online ASR mode:

Set the `use_oline_asr` parameter to `True`.

```
# yahboom.yaml

asr:
  ros__parameters:
    use_online_asr: True # Critical: Change from False to True
```

3.2 Startup and Testing

1. Startup Command:

```
ros2 launch largemodel asr_only.launch.py
```

2. Test:

Say to the microphone: Hello yahboom, it will respond with "I'm here", and then you can start speaking. Finally, it will display the recorded audio converted into text printed on the terminal.

```
sunrise@ubuntu:~/yahboom_ws$ ros2 launch largemodel asr_only.launch.py
[INFO] [launch]: All log files can be found below /home/sunrise/.ros/log/2025-08-07-15-40-45-012409-ubuntu-7665
[INFO] [launch]: Default logging verbosity is set to INFO
[INFO] [asr-1]: process started with pid [7672]
[asr-1] [INFO] [1754552481.072012346] [asr]: The asr model :SenseVoiceSmall is loaded
[asr-1] [INFO] [1754552481.517297362] [asr]: asr_node Initialization completed
[asr-1] [INFO] [1754552491.147085778] [asr]: I'm here
```